

# CMPSC 465: LECTURE XII


## DFS on Directed Graphs

Ke Chen

September 26, 2025

## Recall from the previous lecture ...


On an undirected graph:



## Recall from the previous lecture ...

On an undirected graph:



- Explore a node:



## Recall from the previous lecture ...

On an **undirected** graph:



- ▶ Explore a node:
  - ▶ reveals one connected component



# Recall from the previous lecture ...

On an undirected graph:


- ▶ Explore a node:
  - ▶ reveals one connected component
  - ▶ produces a DFS tree  $\rightsquigarrow$  tree edges, back edges



## Recall from the previous lecture ...

On an **undirected** graph:



- ▶ Explore a node:
  - ▶ reveals one connected component
  - ▶ produces a DFS tree  $\rightsquigarrow$  **tree edges**, **back edges**
- ▶ DFS repeatedly calls Explore:



## Recall from the previous lecture ...

On an **undirected** graph:



- ▶ Explore a node:
  - ▶ reveals one connected component
  - ▶ produces a DFS tree  $\rightsquigarrow$  **tree edges**, **back edges**
- ▶ DFS repeatedly calls Explore:
  - ▶ finds all connected components



## Recall from the previous lecture ...

On an **undirected** graph:

- ▶ Explore a node:
  - ▶ reveals one connected component
  - ▶ produces a DFS tree  $\rightsquigarrow$  **tree edges**, **back edges**
- ▶ DFS repeatedly calls Explore:
  - ▶ finds all connected components  $\rightsquigarrow$  "Can I go from A to B?"







## Recall from the previous lecture ...

On an **undirected** graph:


- ▶ Explore a node:
  - ▶ reveals one connected component
  - ▶ produces a DFS tree  $\rightsquigarrow$  **tree edges**, **back edges**
- ▶ DFS repeatedly calls Explore:
  - ▶ finds all connected components  $\rightsquigarrow$  "Can I go from A to B?"
  - ▶ yields a DFS forest



## Recall from the previous lecture ...

On an **undirected** graph:

- ▶ Explore a node:
  - ▶ reveals one connected component
  - ▶ produces a DFS tree  $\rightsquigarrow$  **tree edges**, **back edges**
- ▶ DFS repeatedly calls Explore:
  - ▶ finds all connected components  $\rightsquigarrow$  "Can I go from A to B?"
  - ▶ yields a DFS forest  $\rightsquigarrow$  **cycle detection**



## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

```
// pre and post are integer arrays of size  $|V|$   
// clock is a global integer counter starting at 1  
Explore( $G, s, color$ )
```

```
    visited[ $s$ ] = color
```

```
    pre[ $s$ ] = clock
```

```
    clock = clock + 1
```

```
    foreach edge  $\{s, v\} \in E$  do
```

```
        if visited[ $v$ ] == 0 then
```

```
            Explore( $G, v, color$ )
```

```
    post[ $s$ ] = clock
```


```
    clock = clock + 1
```

## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

### Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example






## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example






## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example






## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example






## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example





## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example





## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

### Example





## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example







## Pre- and post-visit timestamps (undirected)

We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example




## Pre- and post-visit timestamps (undirected)



We can collect more information during Explore by keeping a global clock that ticks every time we:

- ▶ visit a node for the first time, and
- ▶ leave a node for good.

Example





## DFS in directed graphs





## DFS in directed graphs

Exactly the same algorithm works!



## DFS in directed graphs



Exactly the same algorithm works!



## DFS in directed graphs

Exactly the same algorithm works!

But we have more edge types:





# DFS in directed graphs

Exactly the same algorithm works!

But we have more edge types:

- ▶ **forward edges** lead to a non-child descendant.





# DFS in directed graphs

Exactly the same algorithm works!

But we have more edge types:

- ▶ **forward edges** lead to a non-child descendant.
- ▶ **back edges** lead to an ancestor.







# DFS in directed graphs

Exactly the same algorithm works!

But we have more edge types:


- ▶ **forward edges** lead to a non-child descendant.
- ▶ **back edges** lead to an ancestor.
- ▶ **cross edges** lead to neither an ancestor nor a descendant.



## Pre-, post-visit numbers and edge types

**Fact** If vertex  $w$  is an ancestor of vertex  $v$  in the DFS tree, then

$$pre[w] < pre[v] < post[v] < post[w]$$




## Pre-, post-visit numbers and edge types

**Fact** If vertex  $w$  is an ancestor of vertex  $v$  in the DFS tree, then

$$pre[w] < pre[v] < post[v] < post[w]$$

For an edge  $(w, v)$  in the graph:

- ▶ if  $pre[w] < pre[v] < post[v] < post[w]$ , then  $(w, v)$  is a tree or **forward edge**.




## Pre-, post-visit numbers and edge types

**Fact** If vertex  $w$  is an ancestor of vertex  $v$  in the DFS tree, then

$$pre[w] < pre[v] < post[v] < post[w]$$

For an edge  $(w, v)$  in the graph:

- ▶ if  $pre[w] < pre[v] < post[v] < post[w]$ , then  $(w, v)$  is a tree or **forward edge**.
- ▶ if  $pre[v] < pre[w] < post[w] < post[v]$ , then  $(w, v)$  is a **back edge**.




# Pre-, post-visit numbers and edge types

**Fact** If vertex  $w$  is an ancestor of vertex  $v$  in the DFS tree, then

$$pre[w] < pre[v] < post[v] < post[w]$$

For an edge  $(w, v)$  in the graph:

- ▶ if  $pre[w] < pre[v] < post[v] < post[w]$ , then  $(w, v)$  is a tree or **forward edge**.
- ▶ if  $pre[v] < pre[w] < post[w] < post[v]$ , then  $(w, v)$  is a **back edge**.
- ▶ if  $pre[v] < post[v] < pre[w] < post[w]$ , then  $(w, v)$  is a **cross edge**.



## Cycle detection (directed)

**Fact** A directed graph has a cycle if and only if its DFS forest has a back edge.

## Cycle detection (directed)

**Fact** A directed graph has a cycle if and only if its DFS forest has a back edge.

### Algorithm for cycle detection

1. Run DFS and assign *pre* and *post* numbers.
2. Iterate through all edges  $(w, v)$  and check if  $pre[v] < pre[w] < post[w] < post[v]$ .
3. If found, output “found cycle”; otherwise return “no cycle”.

## Cycle detection (directed)

**Fact** A directed graph has a cycle if and only if its DFS forest has a back edge.

### Algorithm for cycle detection

1. Run DFS and assign *pre* and *post* numbers.
2. Iterate through all edges  $(w, v)$  and check if  $pre[v] < pre[w] < post[w] < post[v]$ .
3. If found, output “found cycle”; otherwise return “no cycle”.

Time complexity?



## Topological sort

A directed graph without cycles is called a **DAG** (directed acyclic graph).

## Topological sort

A directed graph without cycles is called a **DAG** (directed acyclic graph).


DAGs are very useful and are much easier to work with, because we can order the vertices such that all edges point from an earlier vertex to a later vertex.

## Topological sort

A directed graph without cycles is called a **DAG** (directed acyclic graph).

DAGs are very useful and are much easier to work with, because we can order the vertices such that all edges point from an earlier vertex to a later vertex.

Example




## Topological sort

A directed graph without cycles is called a **DAG** (directed acyclic graph).

DAGs are very useful and are much easier to work with, because we can order the vertices such that all edges point from an earlier vertex to a later vertex.

### Example




## Topological sort

A directed graph without cycles is called a **DAG** (directed acyclic graph).

DAGs are very useful and are much easier to work with, because we can order the vertices such that all edges point from an earlier vertex to a later vertex.

### Example



Such an ordering is called a **topological order** or a **linearization** of the DAG.

# Topological sort

**Fact** All DAGs have at least one topological order.

# Topological sort

**Fact** All DAGs have at least one topological order.

**Fact** In a DAG, every edge leads to a vertex with a smaller post number.

# Topological sort

**Fact** All DAGs have at least one topological order.

**Fact** In a DAG, every edge leads to a vertex with a smaller post number.

*Proof:* DAGs do not have back edges, all other edge types satisfies this.



# Topological sort

**Fact** All DAGs have at least one topological order.

**Fact** In a DAG, every edge leads to a vertex with a smaller post number.

*Proof:* DAGs do not have back edges, all other edge types satisfies this.

**Algorithm for topological sort**

1. Run DFS and assign *pre* and *post* numbers.
2. Sort vertices in decreasing order of post numbers.

# Topological sort

**Fact** All DAGs have at least one topological order.

**Fact** In a DAG, every edge leads to a vertex with a smaller post number.

*Proof:* DAGs do not have back edges, all other edge types satisfies this.

**Algorithm for topological sort**

1. Run DFS and assign *pre* and *post* numbers.
2. Sort vertices in decreasing order of post numbers.

**Time complexity?**

# Topological sort

**Fact** All DAGs have at least one topological order.

**Fact** In a DAG, every edge leads to a vertex with a smaller post number.

*Proof:* DAGs do not have back edges, all other edge types satisfies this.

**Algorithm for topological sort**

1. Run DFS and assign *pre* and *post* numbers.
2. Sort vertices in decreasing order of post numbers.

**Time complexity?**  $O(|V| \log |V| + |E|)$ .

# Topological sort

**Fact** All DAGs have at least one topological order.

**Fact** In a DAG, every edge leads to a vertex with a smaller post number.

*Proof:* DAGs do not have back edges, all other edge types satisfies this.

## Algorithm for topological sort

1. Run DFS and assign *pre* and *post* numbers.
2. Sort vertices in decreasing order of post numbers.

**Time complexity?**  $O(|V| \log |V| + |E|)$ .

**Can we do better?**

## Topological sort (faster)

**Idea** Sort at the same time as we assign post numbers.

## Topological sort (faster)

**Idea** Sort at the same time as we assign post numbers.

```
// top_order will contain the ordered vertices
```

```
// idx = |V|
```

```
Explore( $G, s, color$ )
```

```
     $visited[s] = color$ 
```

```
     $pre[s] = clock, \quad clock = clock + 1$ 
```

```
    foreach edge  $\{s, v\} \in E$  do
```

```
        if  $visited[v] == 0$  then
```

```
            Explore( $G, v, color$ )
```

```
     $post[s] = clock, \quad clock = clock + 1$ 
```

```
     $top\_order[idx] = s, \quad idx = idx - 1$ 
```

## Topological sort (faster)

**Idea** Sort at the same time as we assign post numbers.

```
// top_order will contain the ordered vertices
```

```
// idx = |V|
```

```
Explore( $G, s, color$ )
```

```
     $visited[s] = color$ 
```

```
     $pre[s] = clock, \quad clock = clock + 1$ 
```

```
    foreach edge  $\{s, v\} \in E$  do
```

```
        if  $visited[v] == 0$  then
```

```
            Explore( $G, v, color$ )
```

```
     $post[s] = clock, \quad clock = clock + 1$ 
```

```
     $top\_order[idx] = s, \quad idx = idx - 1$ 
```

**Time complexity?**

## Topological sort (faster)

**Idea** Sort at the same time as we assign post numbers.

```
// top_order will contain the ordered vertices
```

```
// idx = |V|
```

```
Explore(G, s, color)
```

```
    visited[s] = color
```

```
    pre[s] = clock,    clock = clock + 1
```

```
    foreach edge  $\{s, v\} \in E$  do
```

```
        if visited[v] == 0 then
```

```
            Explore(G, v, color)
```

```
    post[s] = clock,    clock = clock + 1
```

```
    top_order[idx] = s,    idx = idx - 1
```

**Time complexity?**  $O(|V| + |E|)$ .