



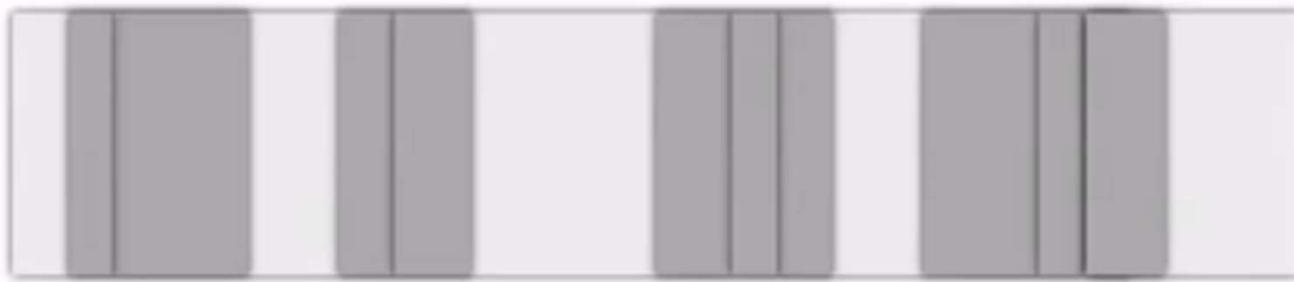
Memory Management

Professor: Suman Saha

Heaps



PennState



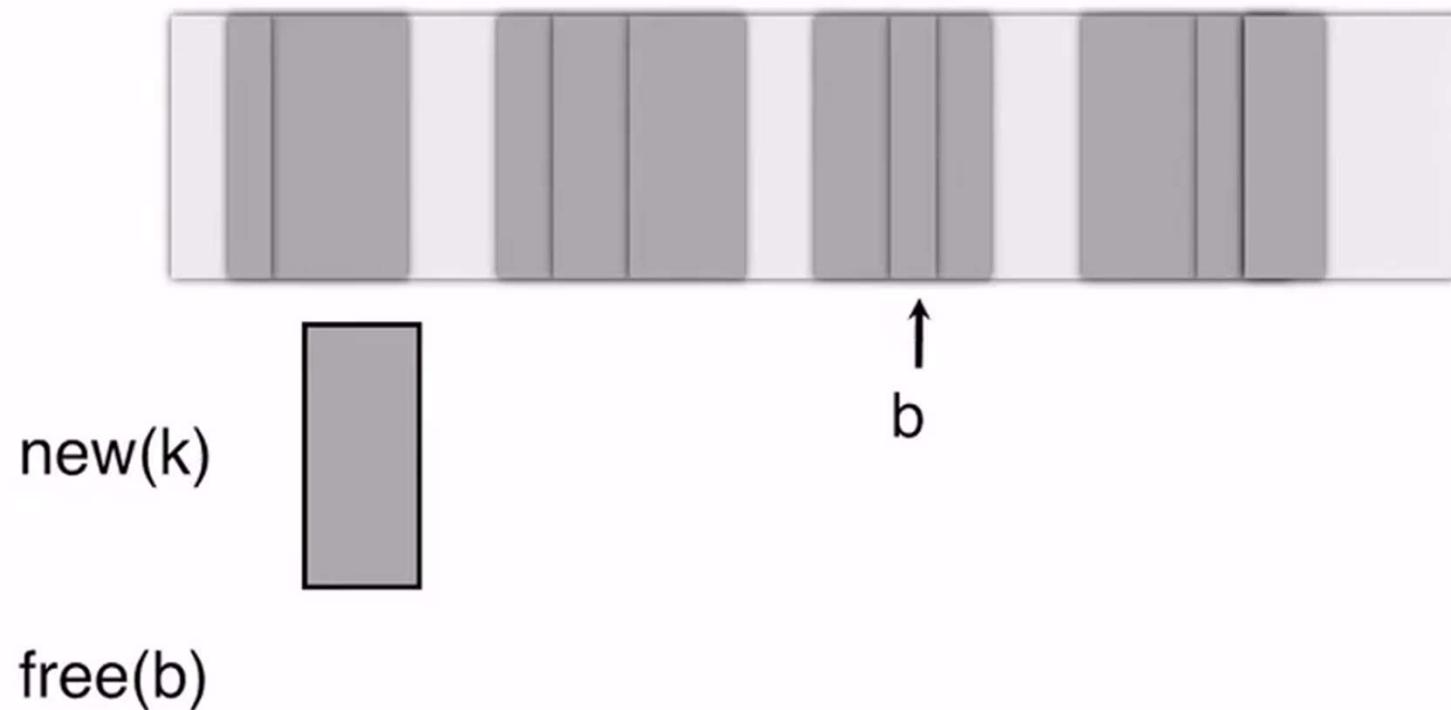
`new(k)`



Heaps



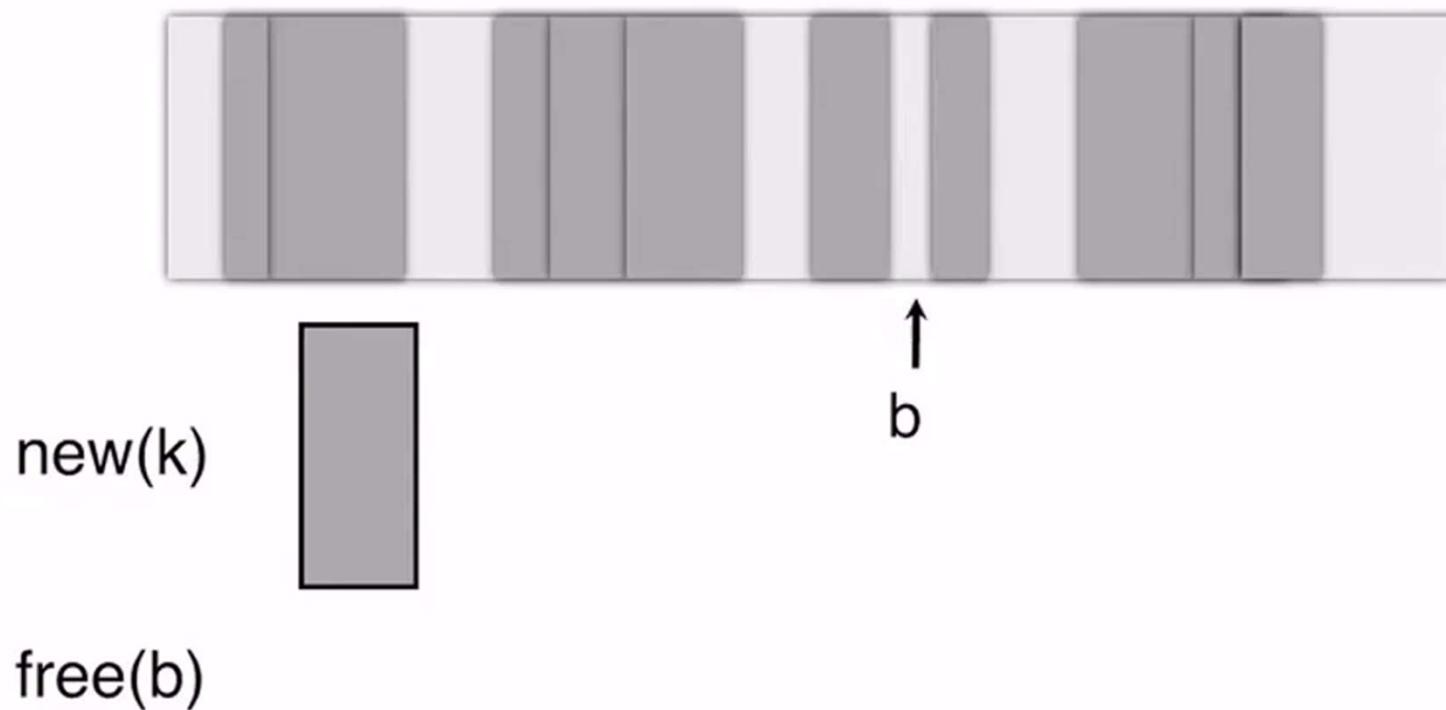
PennState



Heaps



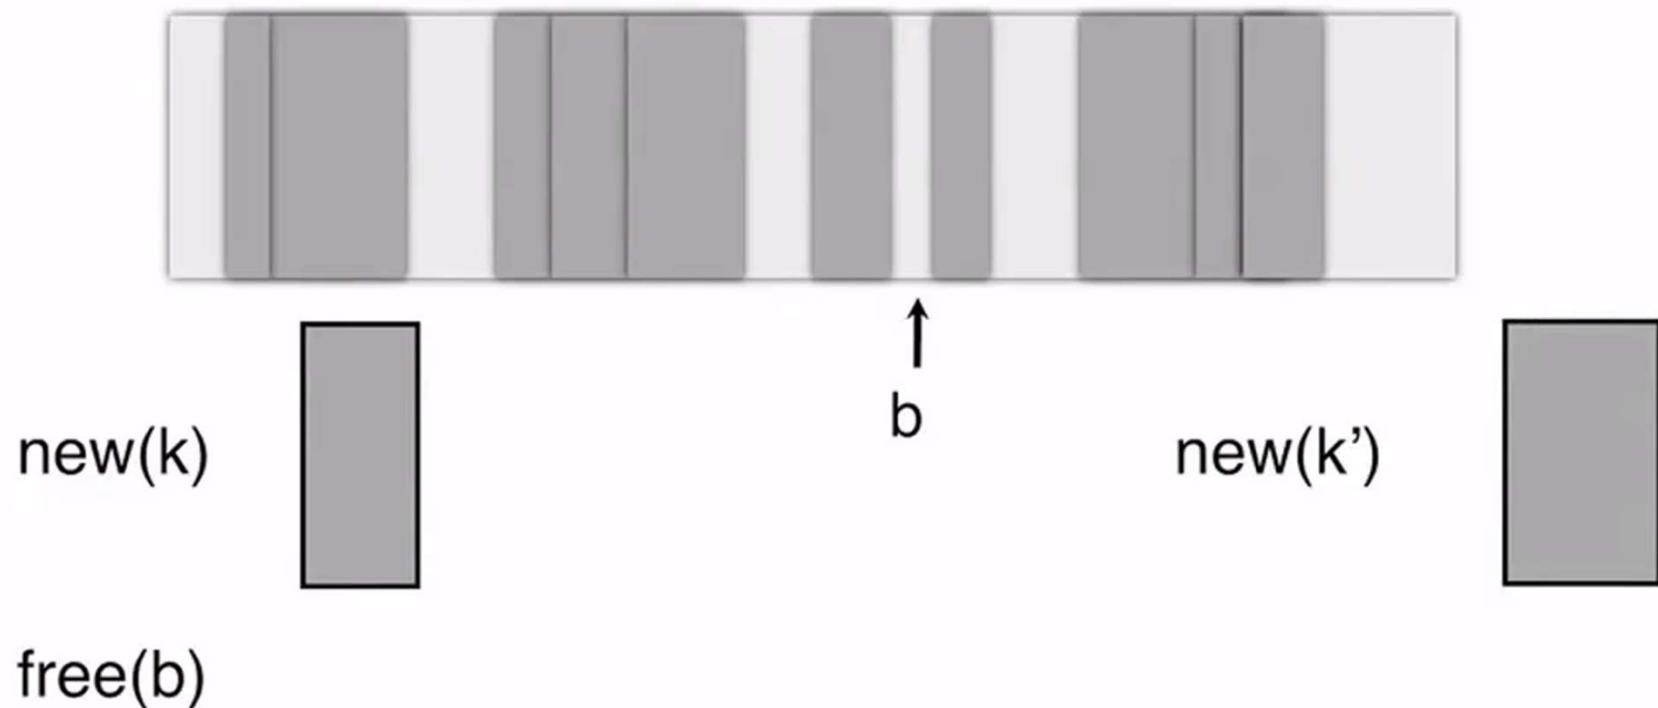
PennState



Heaps



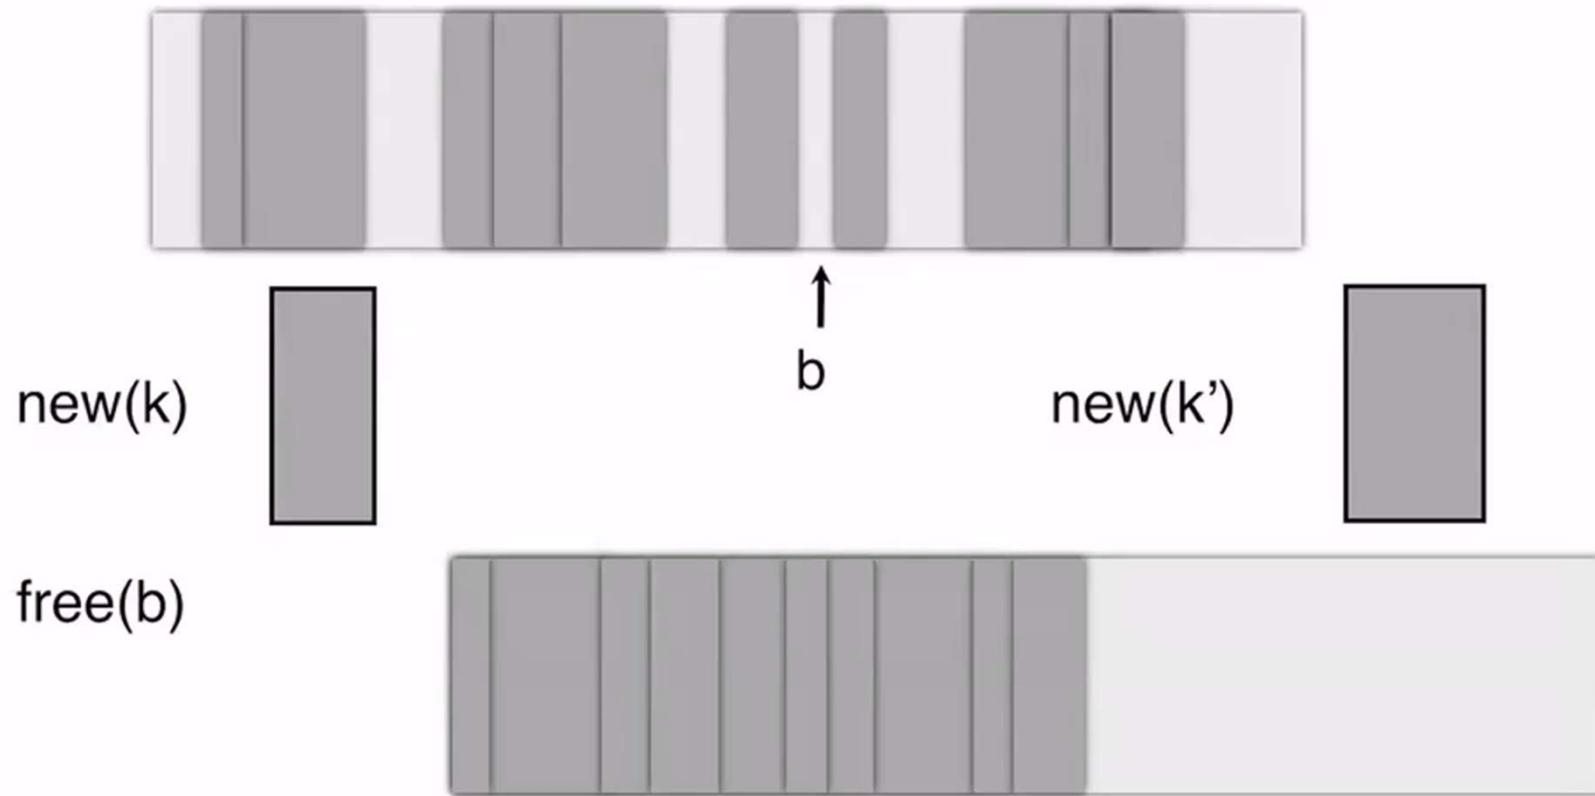
PennState



Heaps



PennState



Heap Management



Allocator: a routine takes size of requested heap space, and search for free space

Usually, heap is managed in blocks. Allocator may return larger block than requested

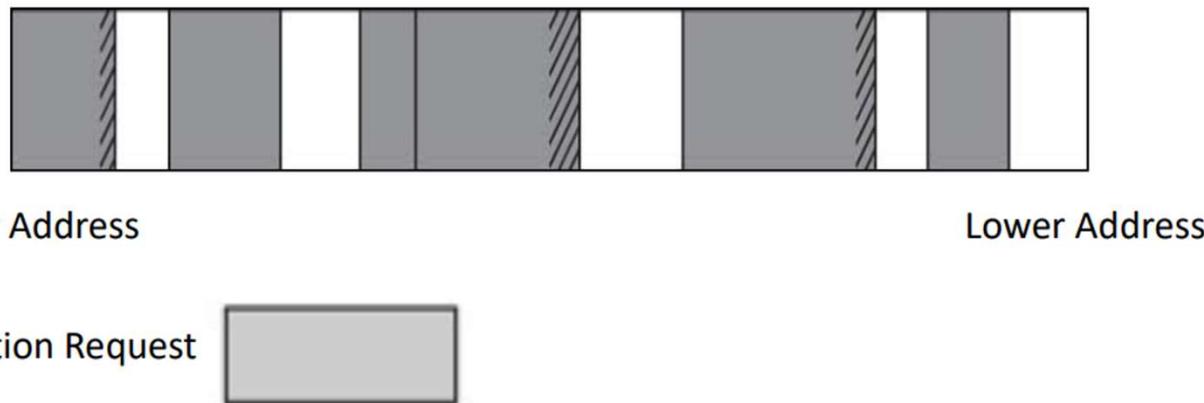
Deallocator: collect free space and merge with other free space when possible

Heap compaction: move all used blocks to one end

Heap-based Allocation



PennState



Internal fragmentation:
the allocation request is smaller than the assigned memory block

External fragmentation:
none of the scattered free space is large enough for the request

Heap Management Algorithms (not covered in this course)



First-fit: select the first free block that is large enough

Best-fit: select the smallest free block that fits

Buddy system: maintain various pools of free blocks with size of 2^k

Fibonacci heap: maintain various pools of free blocks with size of Fibonacci numbers

Heap Management



PennState

Programmer Management (C, C++)

- Pros: implementation simplicity, performance
- Cons: error prone (dangling pointers, memory leaks)

```
Node *p, *q;  
p = new Node();  
q = new Node();  
q = p; // memory leaks  
delete(p); // q becomes dangling pointer
```

Heap Management



Automatic Management (Java, Scheme)

- No dangling pointers, no memory leaks
- Cost: Slower than programmer management

Garbage Collection



Garbage: inaccessible heap objects

```
void foo () {  
    int* a = new int[10];  
    return;  
}
```

Issues of heap management:

- Collect too aggressively: dangling pointers
- Collect too conservatively: memory leaks
- Key problem: collect only objects that are ***inaccessible*** from program



GC I: Reference Counting

Maintain a reference count with each heap object

Set to 1 when object is created

Incremented each time new reference to it is created

Decrement each time reference to it is deleted

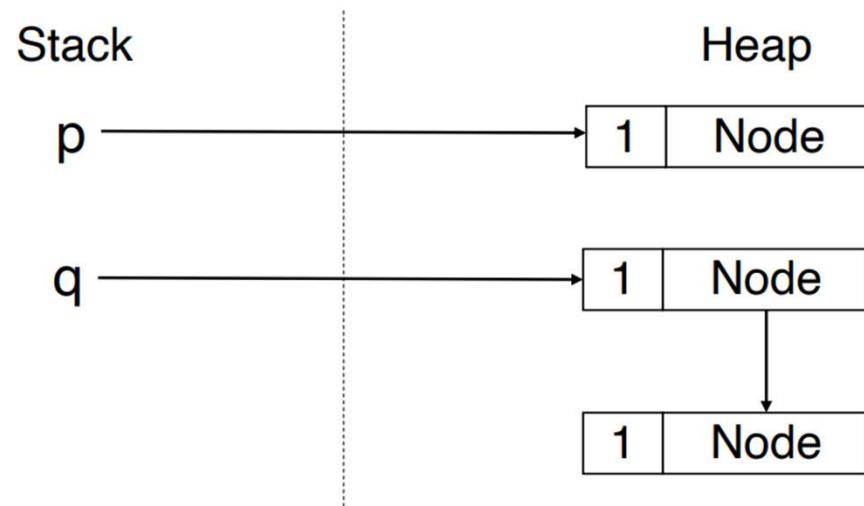
Collect when count becomes 0

GC I: Reference Counting



PennState

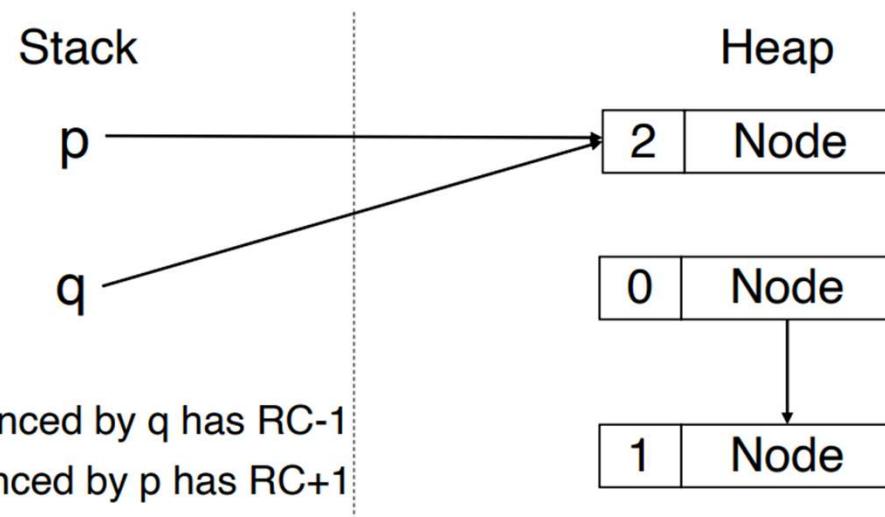
```
Node *p, *q;  
p = new Node();  
q = new Node();  
→ q.next = new Node();  
q = p;  
p = null;
```





GC I: Reference Counting

```
Node *p, *q;  
p = new Node();  
q = new Node();  
q.next = new Node();  
→ q = p;  
p = null;
```

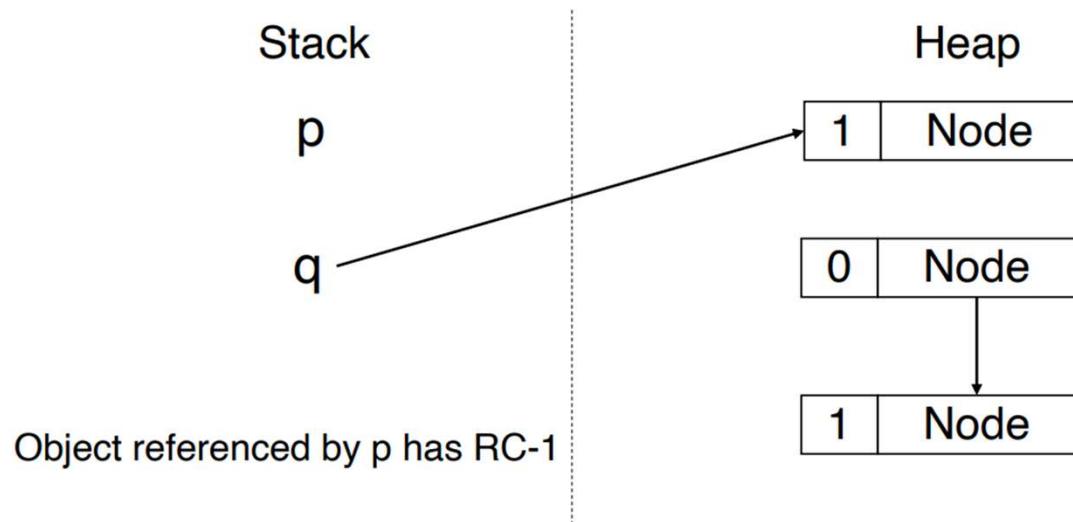


GC I: Reference Counting



PennState

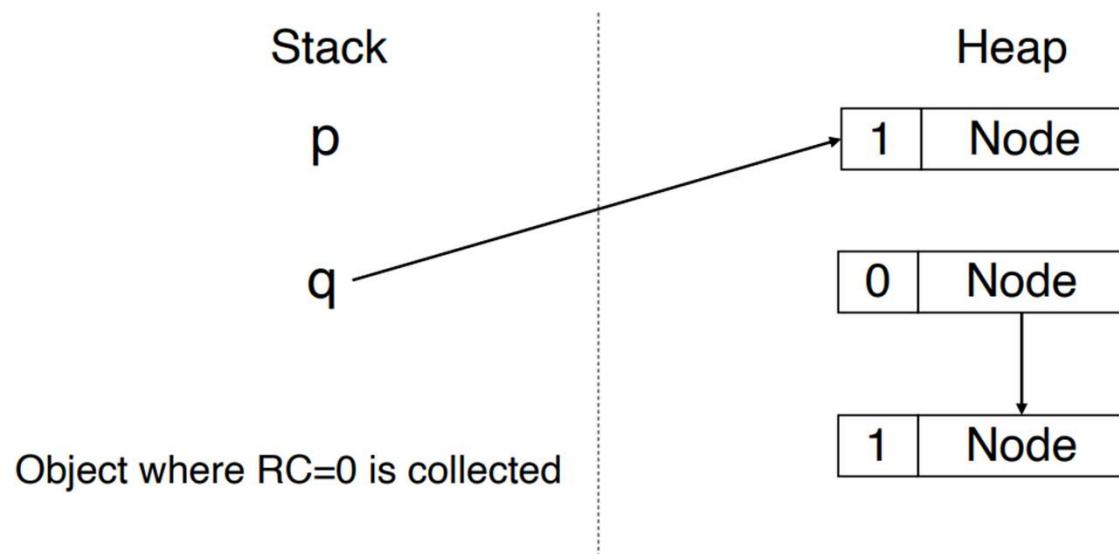
```
Node *p, *q;  
p = new Node();  
q = new Node();  
q.next = new Node();  
q = p;  
p = null;
```





GC I: Reference Counting

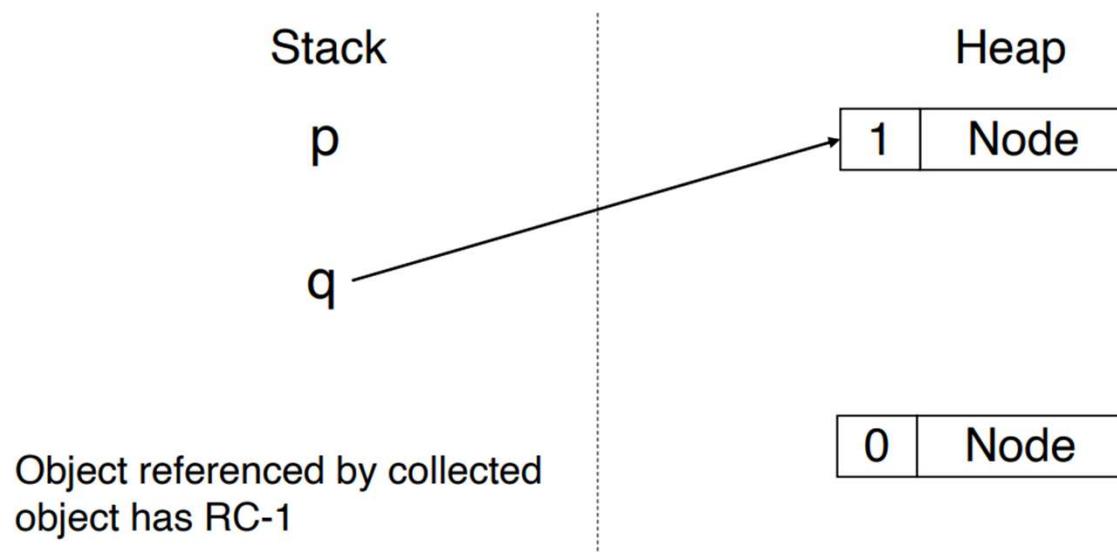
Garbage Collection





GC I: Reference Counting

Garbage Collection





GC I: Reference Counting

When is an object dereferenced?

- Reference is LHS of Assignment
- Reference on stack is destroyed when function returns
- Reference is destroyed when an object with count 0 is collected



GC I: Reference Counting

Reference Counting is about **Object Ownership**

- when one object creates a reference to another object, it owns that object (retain)
- when the object deletes that reference, it relinquishes ownership (release)

Multiple owners of an object

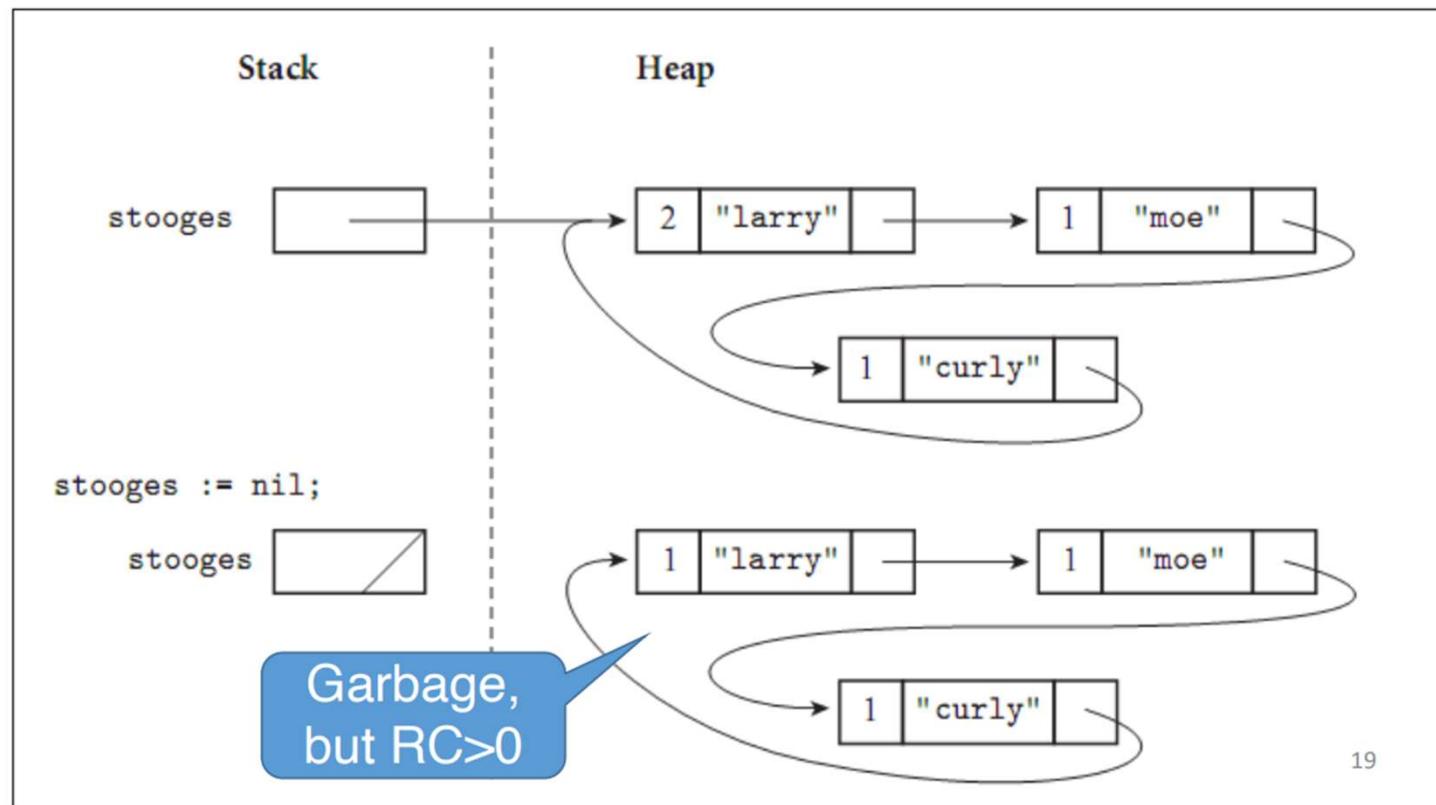
Zero owners of an object

Problem of Reference Counting



PennState

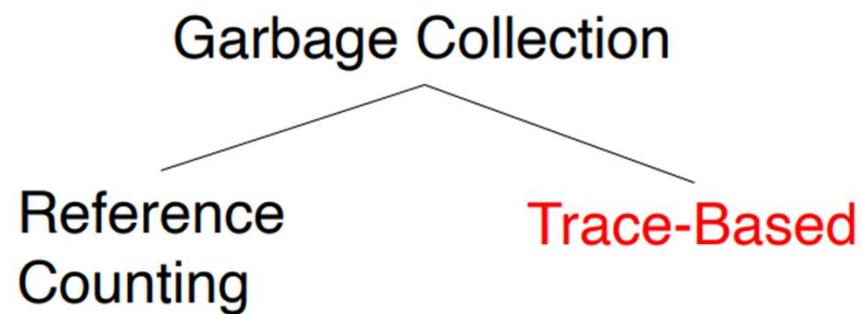
Ownership \neq Accessibility



19



Taxonomy

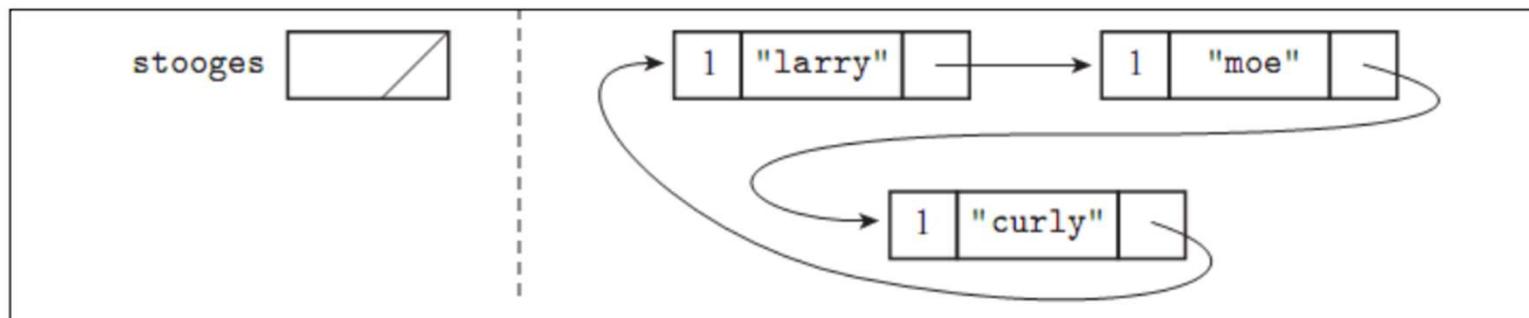




What is Garbage?

Ideally, any heap block not used in the future

In practice, the garbage collector identifies blocks inaccessible from program



Essentially a **reachability problem (from alive variables)**

GC II: Tracing Collection

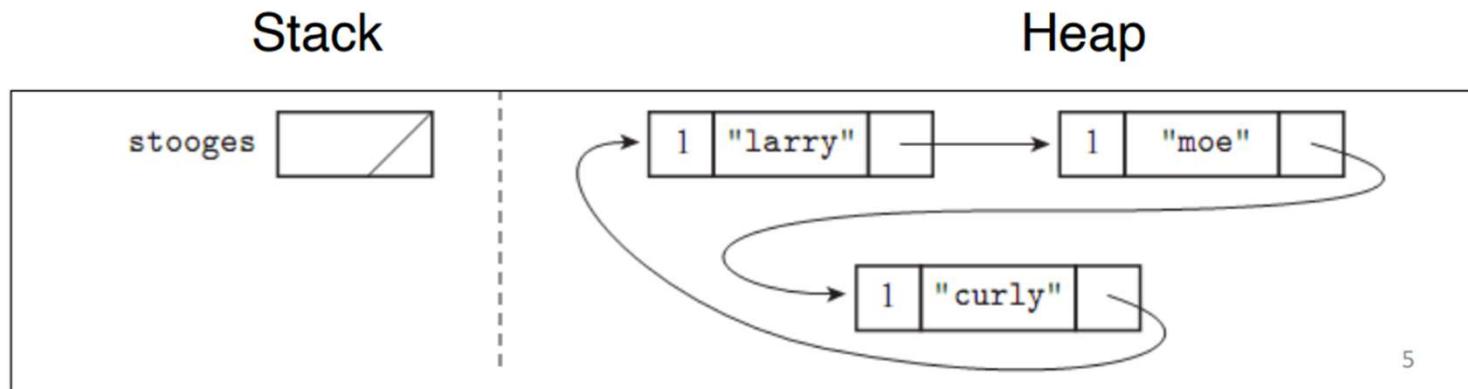


PennState

What allocated blocks are still accessible (live)?

- Reference in registers, static area, or stack
- Reference from reachable objects to other objects

Essentially a reachability problem (heap as directed graph)





Abstraction of Tracing

Root set: references in registers, static area or stack

Heap graph: node → object, edge → pointer/reference

Accessible set = Reachable object from root set

Overview



PennState

Essentially a reachability problem (heap as directed graph)

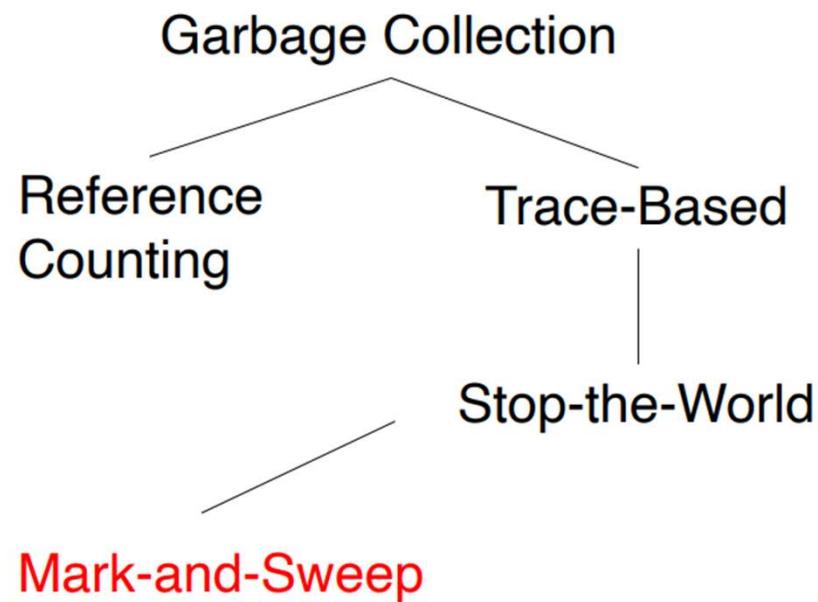
- Start from references in registers/stack/static area
- Traverse the graph to identify reachable blocks
- Remove unreachable blocks

Tradeoffs

- Time overhead - % of time spent on collection
- Space overhead – memory needed for collection algorithm
- Pause time – the time when a program is interrupted
- Frequency of collection
- Promptness – time between a block becomes garbage and it is collected



Taxonomy



Mark-And-Sweep



Triggered when the heap is full, interrupt program

1 Mark Bit (MB) for each object, initially 0

Phase 1 (Mark)

- Traverse graph, set MB to 1 for visited node

Phase 2 (Sweep)

- Traverse **all objects in heap**, collect objects whose MB=0; otherwise, set MB to 0 (prepare for the next collection)

Mark-And-Sweep

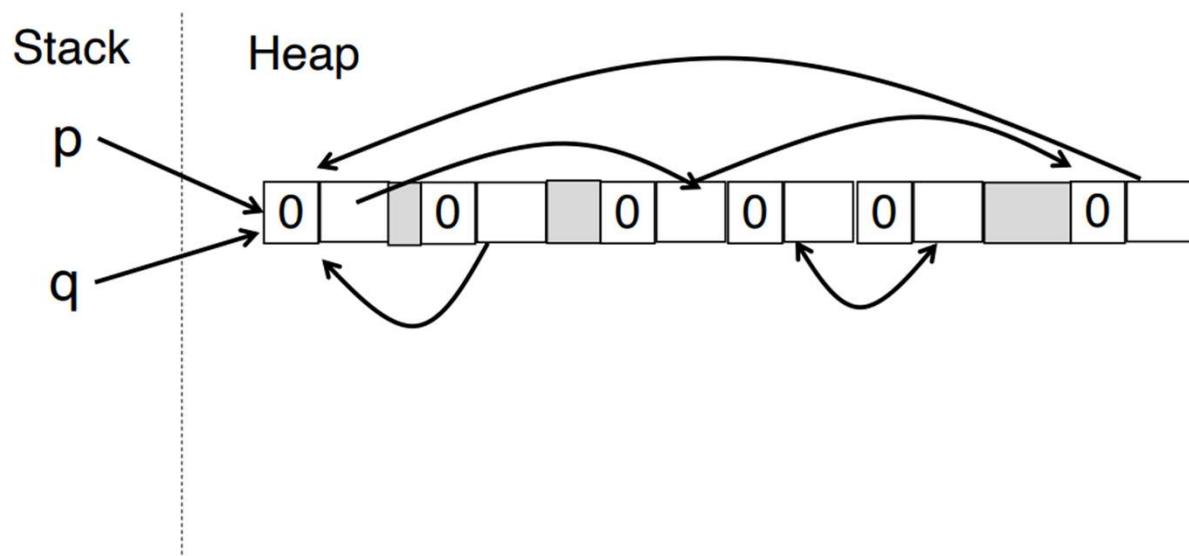


PennState

Example

Initial state

 Free space
 Object with $MB=0$



Mark-And-Sweep

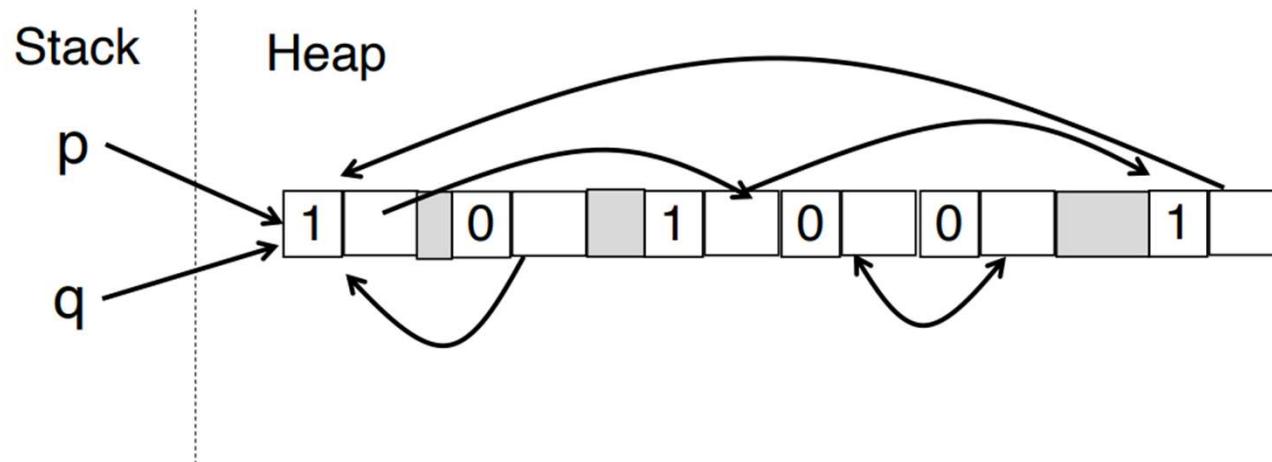


PennState

Example

Mark

 Free space
 Object with $MB=0$



Mark-And-Sweep

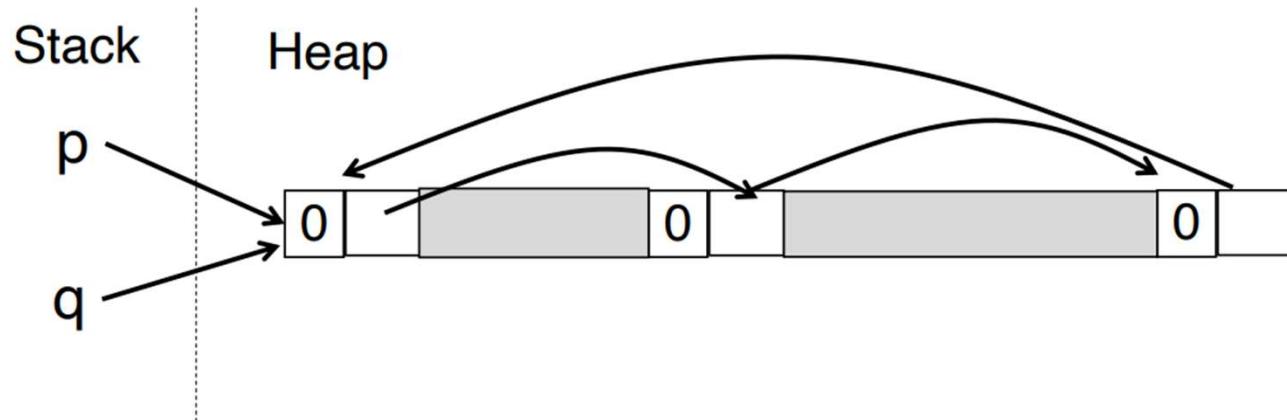


PennState

Example

Sweep

 Free space
 Object with $MB=0$



Mark-And-Sweep



1 Mark Bit (MB) for each object, initially 0

A: number of alive objects; N: all objects on heap

Phase 1 (Mark) – $O(A)$

- Traverse graph, set MB to 1 for visited node

Phase 2 (Sweep) – $O(N)$

- Traverse **entire heap**, collect objects whose MB=0; otherwise, set MB to 0 (prepare for the next collection)



Stop-the-World Collectors

A: number of alive objects; N: all objects on heap

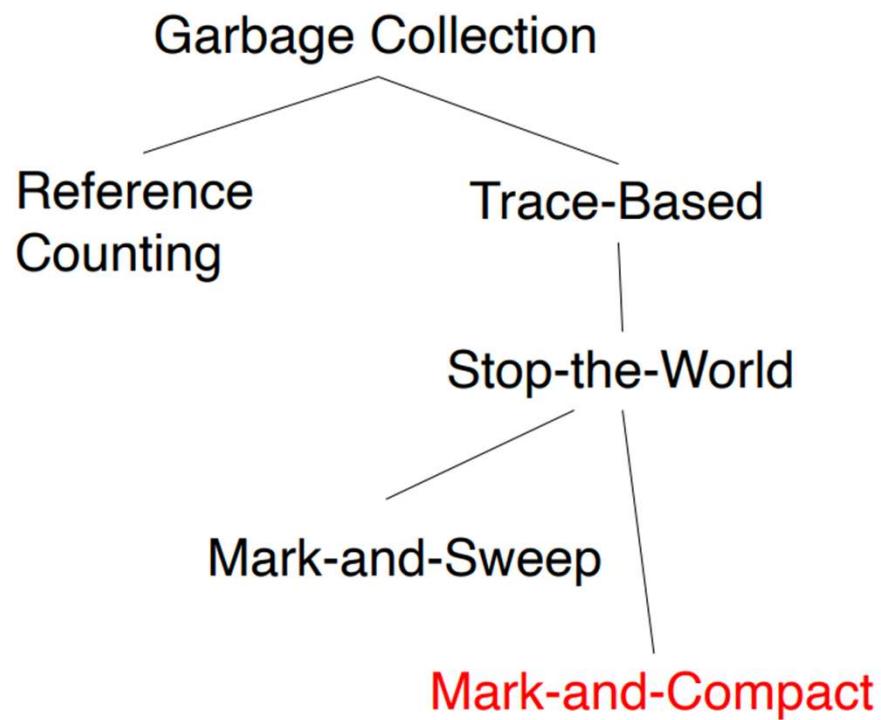
H: total heap size

	Mark-And-Sweep
Complexity	$O(N)$
Fragmentation?	Yes
Memory move	No
Effective heap size	H

Taxonomy



PennState



Mark-and-Compact



PennState

Similar to mark-and-sweep, except
Compact: move reachable object to one end

Mark-and-sweep



Mark-and-compact





Why Compact?

Locality: fewer pages or cache-lines needed to hold the active data.

Reduce fragmentation: available space merged to store large objects



Mark-and-Compact?

Triggered when the heap is full, interrupt program

Phase 1 (Mark): same as mark-and-sweep

Phase 2 (Compact)

a) Maintain **M**, a map from object to new location

I, a pointer initialized to the start of heap

For each heap object **o** from **low address to high address**

if **MB=1**, **M(o)=I**, **I = I + size(o)**

Mark-and-Compact?



Triggered when the heap is full, interrupt program

Phase 1 (Mark): same as mark-and-sweep

Phase 2 (Compact)

- b) For each heap object o from **low address to high address**
 - if $MB=1$, move o to location $M(o)$, update pointers in o
(use the map M), set MB to 0
 - if $MB=0$, collect o
- c) Retarget root references (use M)

Mark-and-Compact?

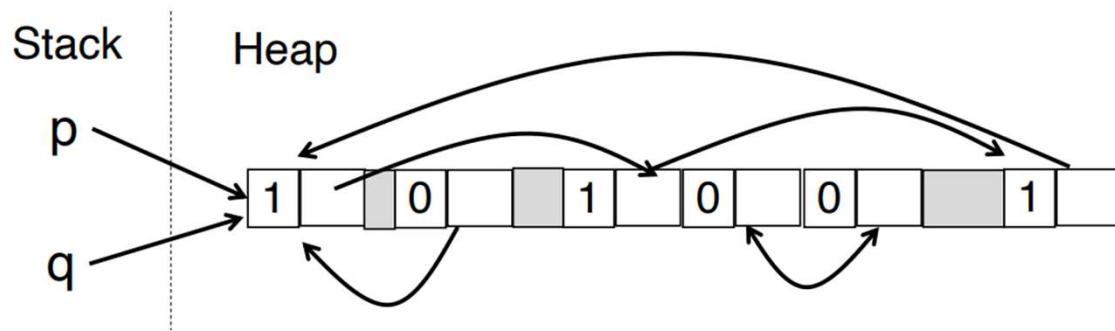


PennState

Example

Mark

Free space
Object with MB=0



Mark-and-Compact?

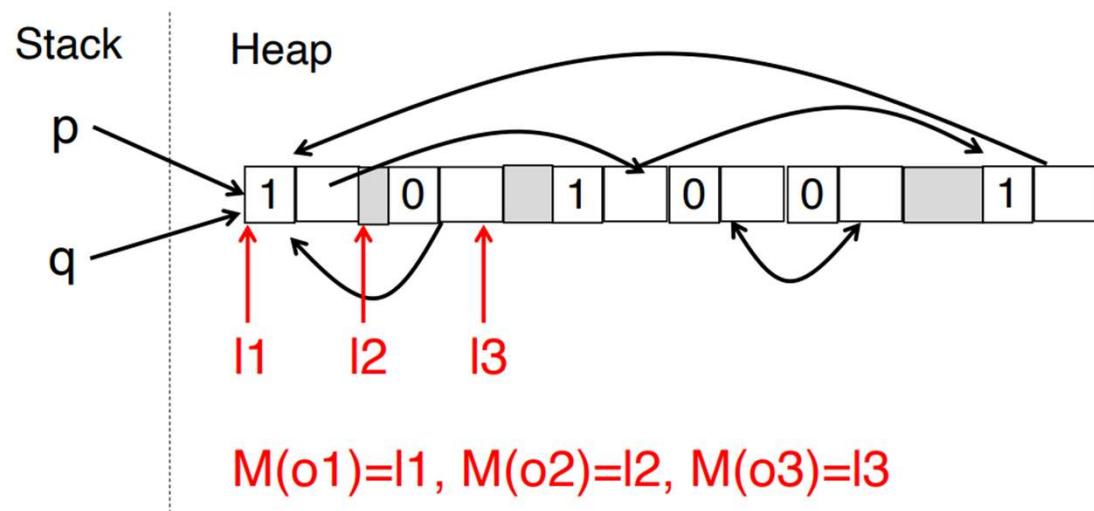


PennState

Example

Compact a

 Free space
 Object with MB=0



Mark-and-Compact?

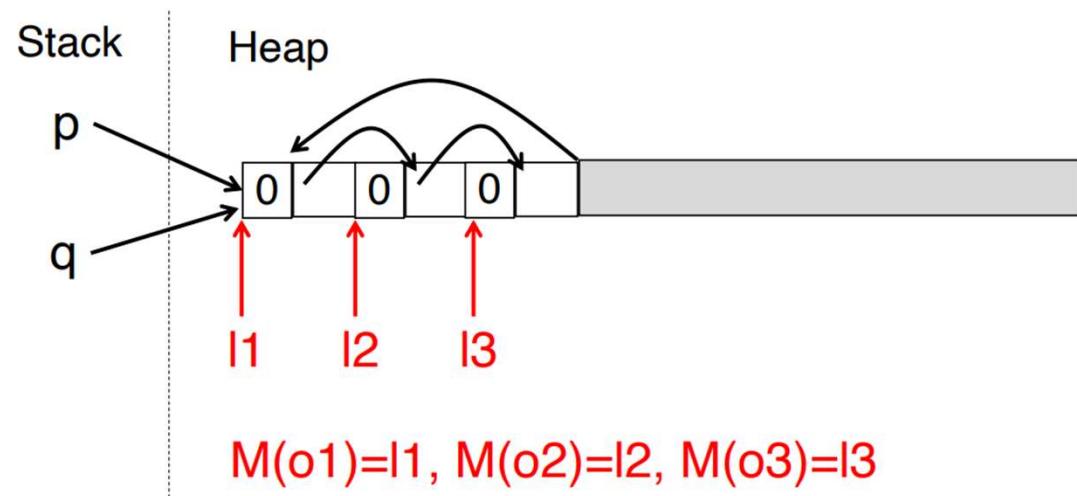


PennState

Example

Compact b&c

 Free space
 Object with MB=0



Mark-and-Compact?



Pros:

Eliminates fragmentation

Great with long-living objects

Cons:

Bad with short-lived data

Collection time is proportional to the heap size

Mark-and-Compact?



A: number of alive objects; N: all objects on heap

Phase 1 (Mark) – $O(A)$

- same as mark-and-sweep

Phase 2 (Compact) – $O(N)$

- Compute new locations for objects
- Move new objects to new location



Stop-the-World Collectors

A: number of alive objects; N: all objects on heap

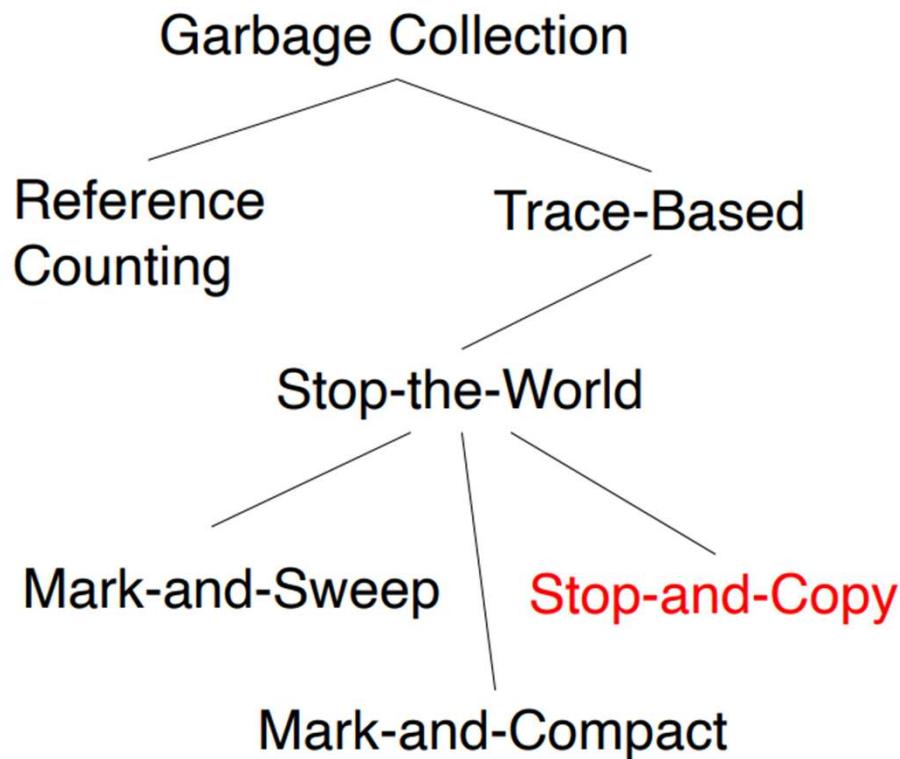
H: total heap size

	Mark-And-Sweep	Mark-And-Compact
Complexity	$O(N)$	$O(N)$
Fragmentation?	Yes	No
Memory move	No	Yes
Effective heap size	H	H

Taxonomy



PennState





Stop-and-Copy (Copying Collector)

Time-Space Tradeoff:

- 2 Heaps: allocate space in one, copy to second when the other is full (half of the heap is unused)
- Only **one pass over live objects** is needed (much faster than previous algorithms)

MB bit is not needed



Stop-and-Copy

Triggered when heap in use is full, interrupt program

For each visited object

1. Copy it to the other heap (no fragmentation)
2. For the object in the old heap, keep a *forwarding pointer* to the new address
3. For each object it points to:
 - a) if visited: update links (follow forwarding pointer)
 - b) otherwise, (recursively) visit object and update link

Retarget root references

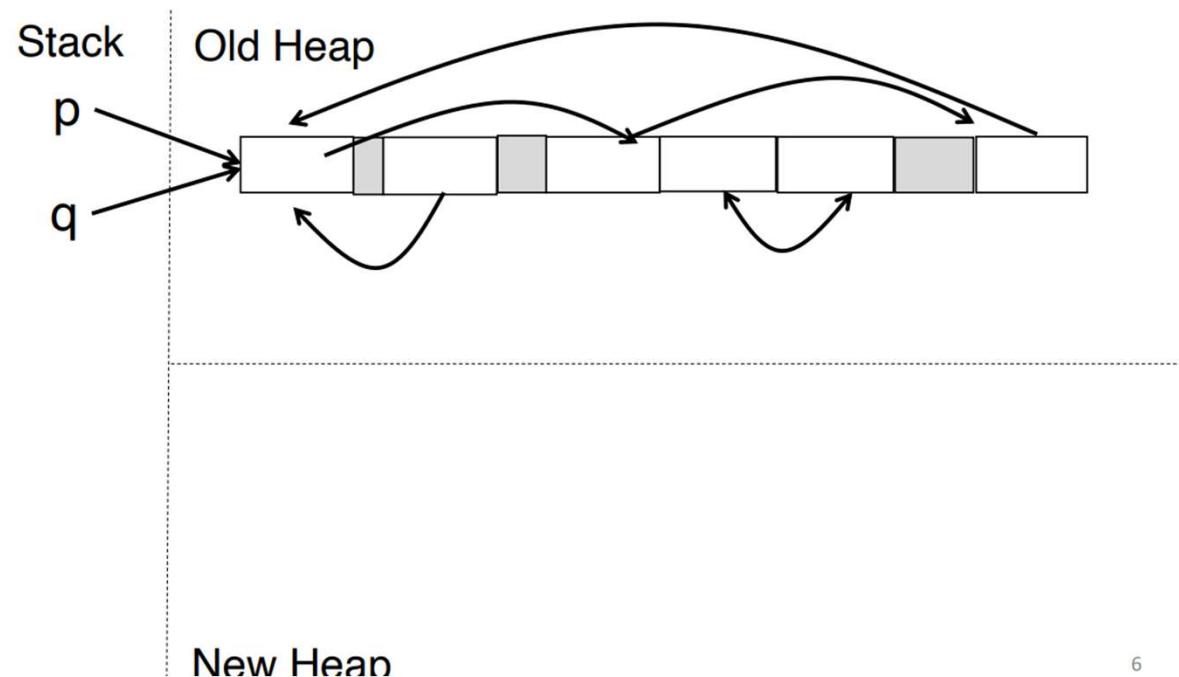
Switch heaps



Stop-and-Copy

Example

Initial state

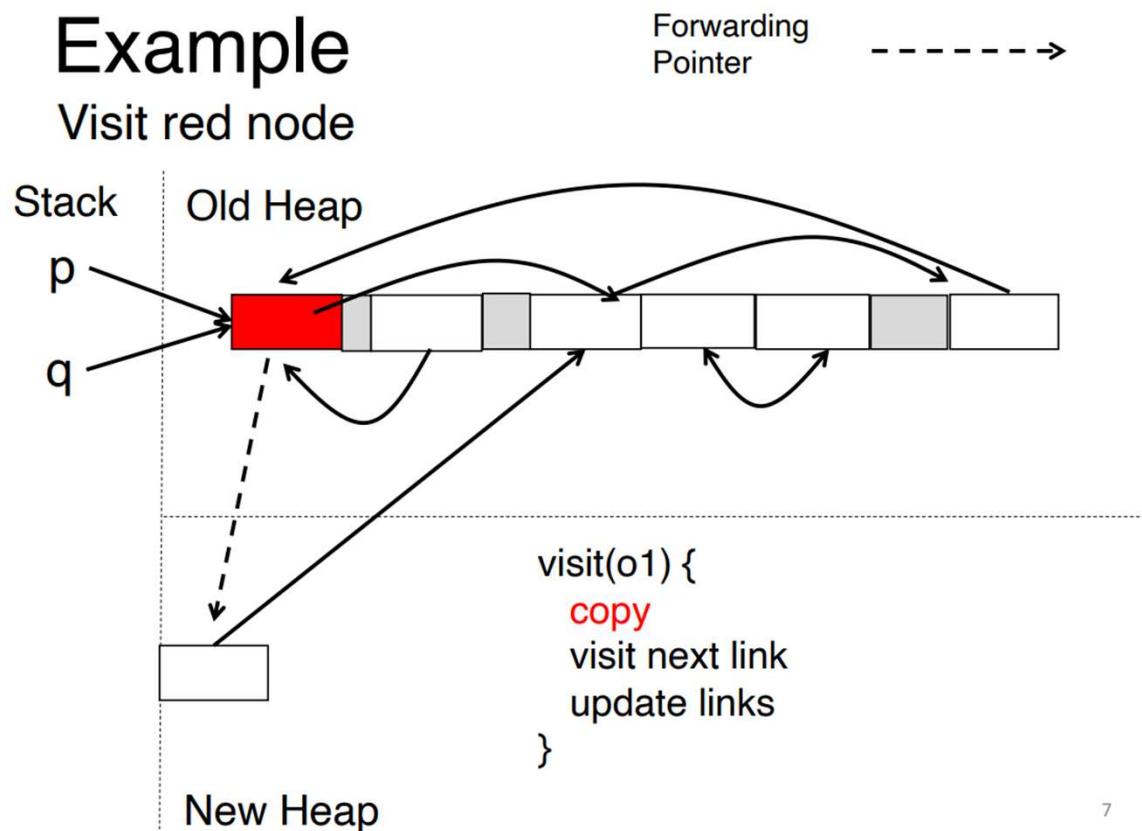




Stop-and-Copy

Example

Visit red node

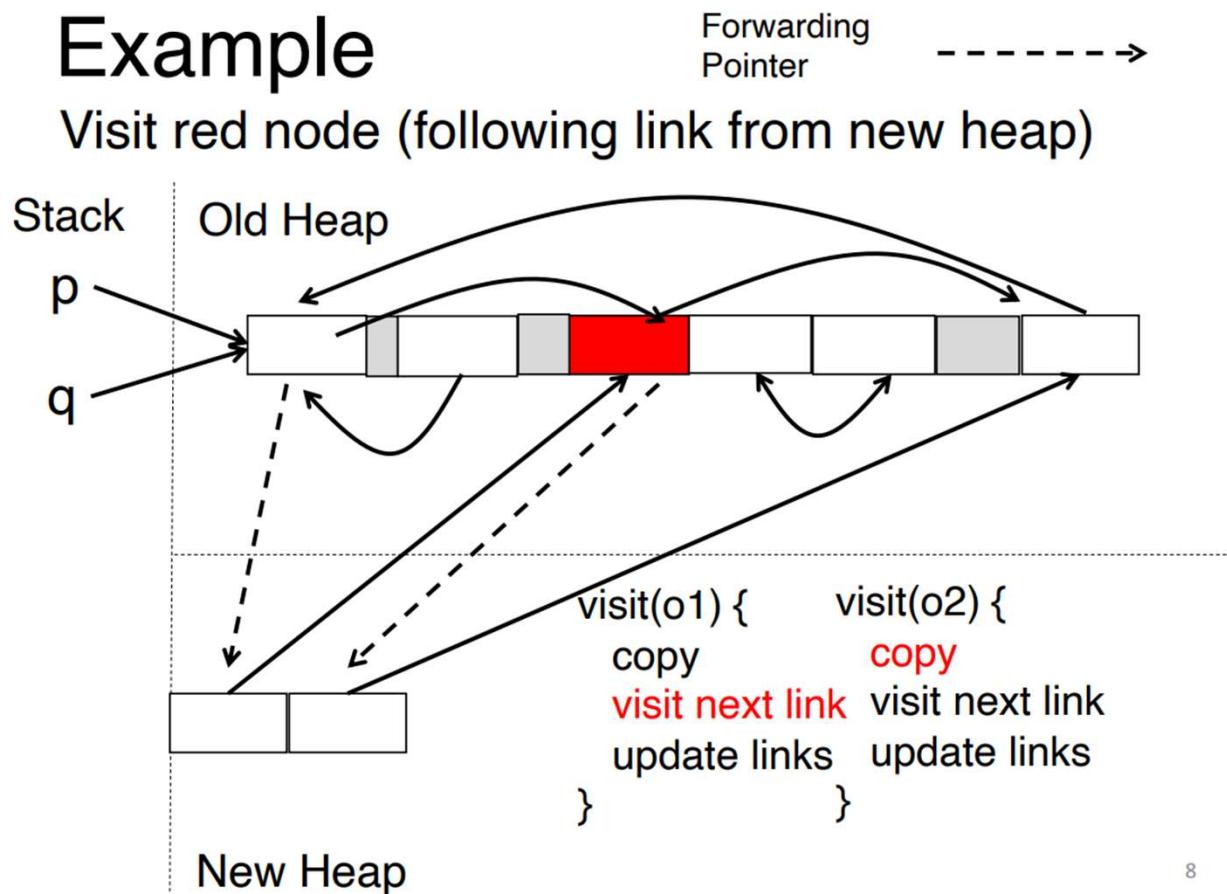




Stop-and-Copy

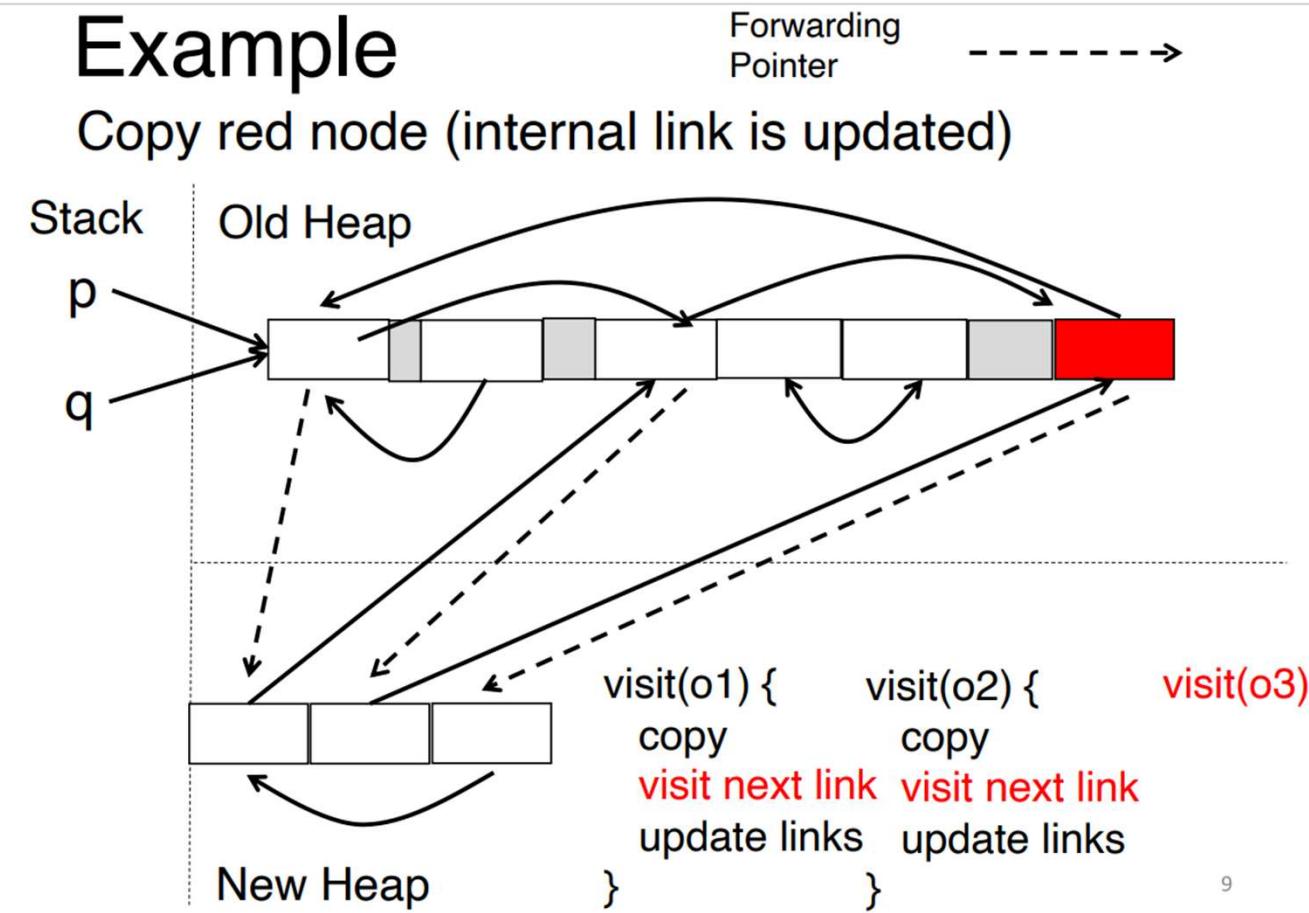
Example

Visit red node (following link from new heap)





Stop-and-Copy

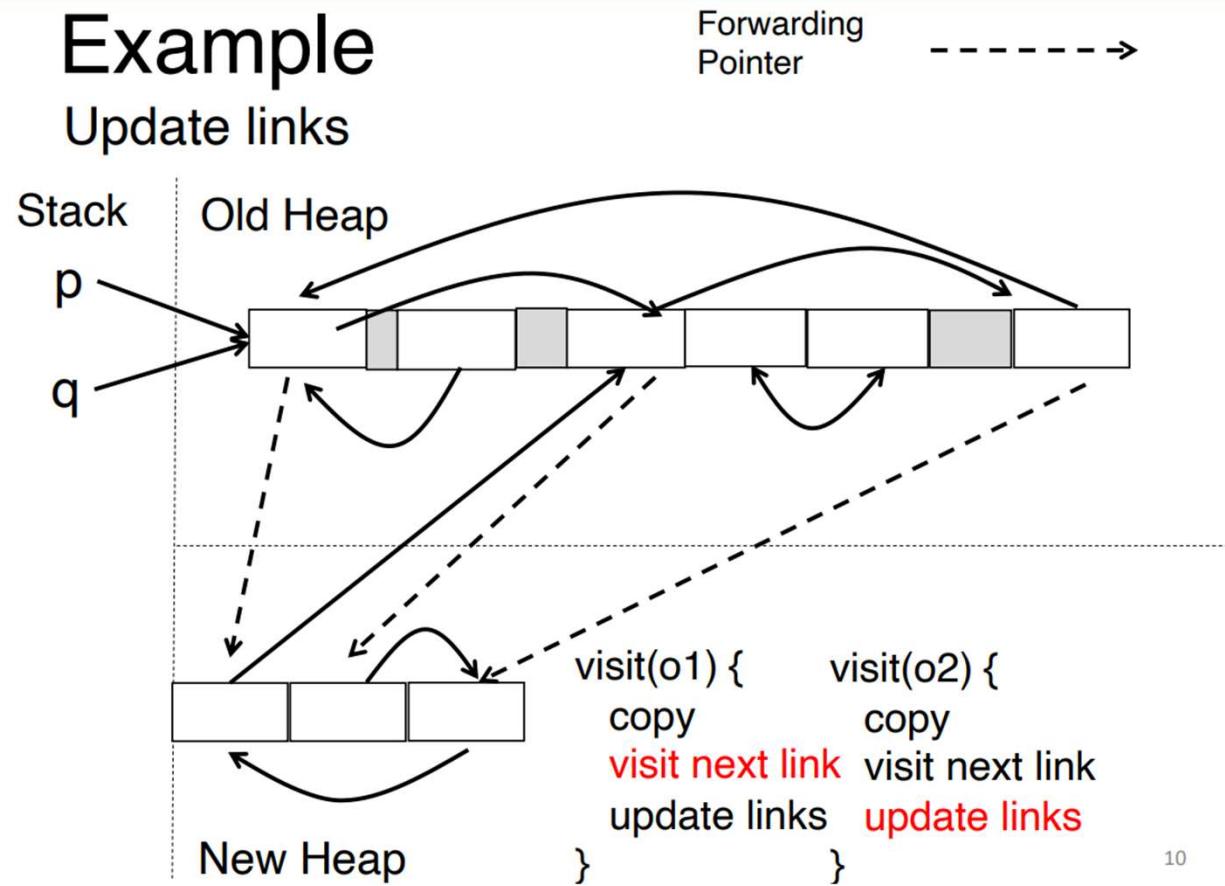




Stop-and-Copy

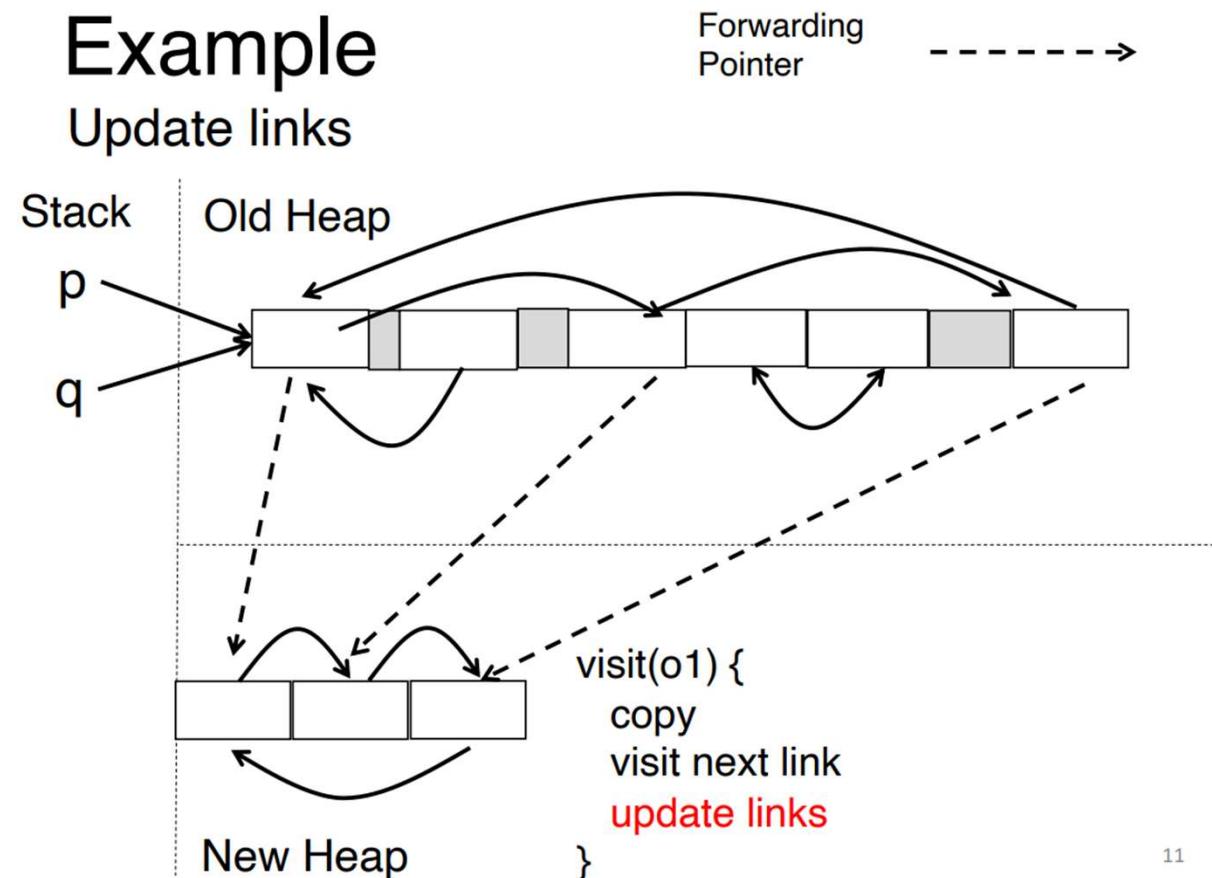
Example

Update links





Stop-and-Copy

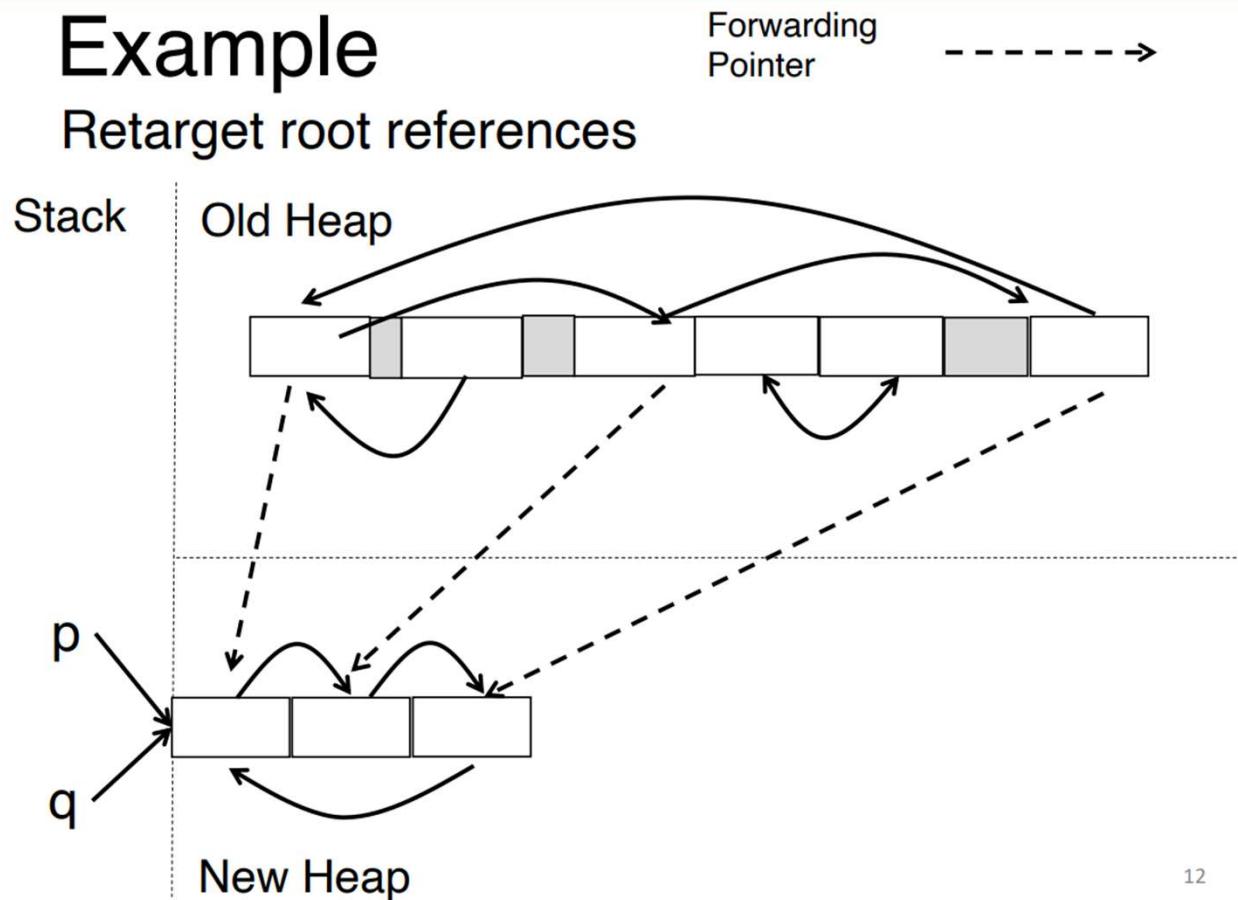




Stop-and-Copy

Example

Retarget root references



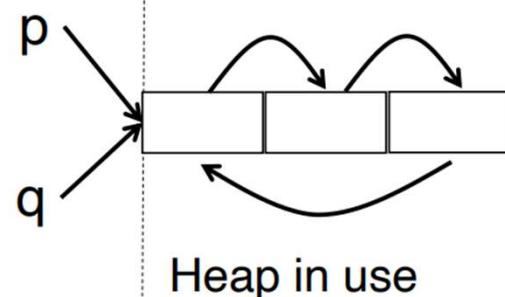


Stop-and-Copy

Example

Swap Heaps (no need to clean the old one)

Stack Heap for copying





Stop-the-World-Collectors

A: number of alive objects; N: all objects on heap

H: total heap size

	Mark-And-Sweep	Mark-And-Compact	Stop-And-Copy
Complexity	$O(N)$	$O(N)$	$O(A)$
Fragmentation?	Yes	No	No
Memory move	No	Yes	Yes
Effective heap size	H	H	$H/2$



Stop-and-Copy

Pros:

Eliminate fragmentation

Collection time proportional to *live* objects

Great for short-lived data

Cons:

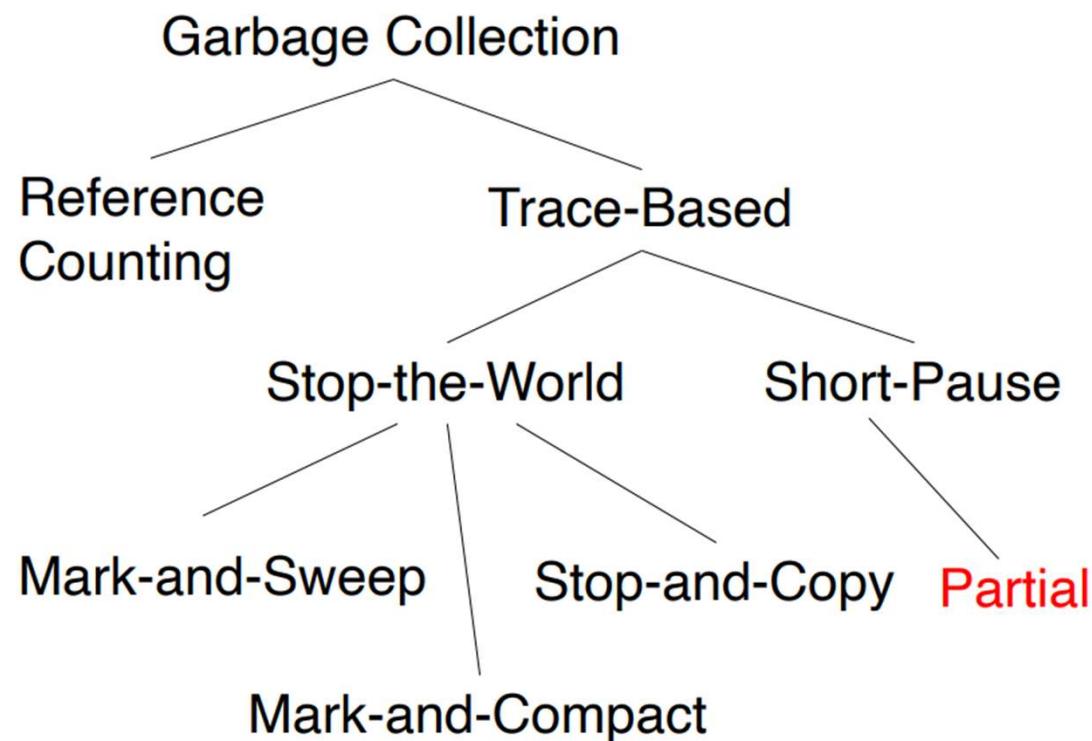
Only half of heap is in use, requires more collecting
(virtual memory alleviates this limitation)

Bad for long living data (copying data back and forth)

Taxonomy

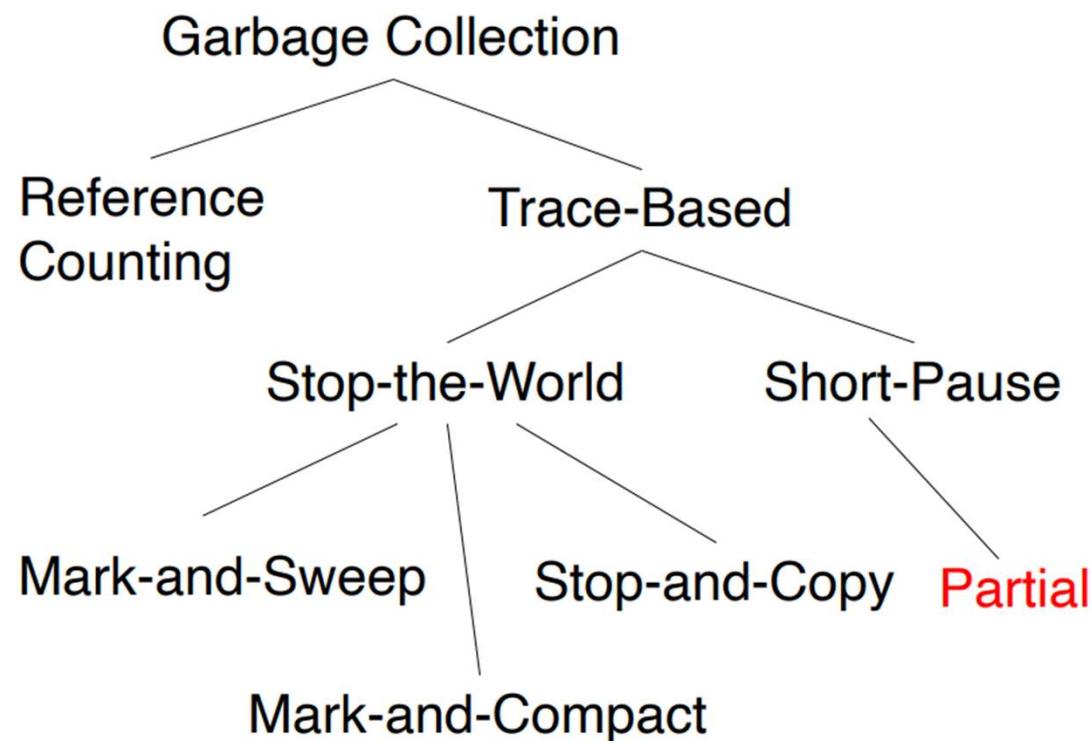


PennState





Short-Pause Collection





PennState

Short-Pause Collection

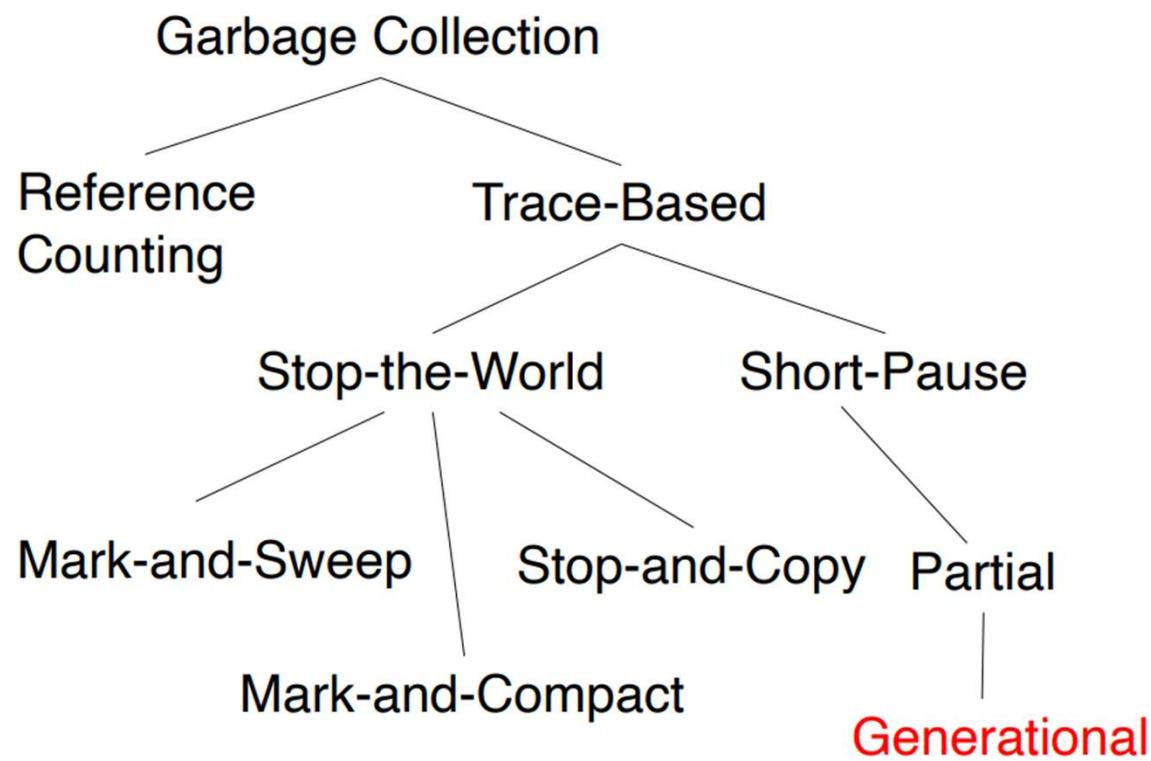
Partial

- stop the program, but only briefly, to garbage collect **a part of the heap**

Taxonomy



PennState





The Object Life-Cycle

“Most objects die young.”

- But those that survive one GC are likely to survive many

Idea: tailor GC to spend more time on regions of the heap where objects have just been created

A better ratio of reclaimed space per unit time

Generational



Divide heap into generations $g_1, g_2 \dots$

- g_i holds older objects than g_{i-1}

Create new objects in g_1 , until it fills up

GC g_1 only; move reachable objects to g_2 after several collections (typically one collection)

Generational



When g_2 fills, garbage collect g_1 and g_2 , and put the reachable objects in g_3

In general: When g_i fills, collect g_1, g_2, \dots, g_i , and put the reachable objects in g_{i+1}

What GC algorithm is better for young generation?
How about old generation?



Algorithm for Young Generation

A: number of alive objects; N: all objects on heap

$A \ll N$

	Mark-And-Sweep	Mark-And-Compact	Stop-And-Copy
Complexity	$O(N)$	$O(N)$	$O(A)$
Fragmentation?	Yes	No	No
Memory move	No	Yes	Yes
Effective heap size	H	H	$H/2$

Stop-and-copy is the most efficient



Algorithm for Old Generation

A: number of alive objects; N: all objects on heap

A is close to N

	Mark-And-Sweep	Mark-And-Compact	Stop-And-Copy
Complexity	$O(N)$	$O(N)$	$O(A)$
Fragmentation?	Yes	No	No
Memory move	No	Yes	Yes
Effective heap size	H	H	$H/2$

Mark-And-Compact removes fragmentation



Algorithm for Old Generation

A: number of alive objects; N: all objects on heap

What if most objects have the same size?

	Mark-And-Sweep	Mark-And-Compact	Stop-And-Copy
Complexity	$O(N)$	$O(N)$	$O(A)$
Fragmentation?	Yes	No	No
Memory move	No	Yes	Yes
Effective heap size	H	H	$H/2$

Mark-And-Sweep saves the cost of moving objects

Generational



Pros:

- Divide heap according to lifetimes of data
- Great for data with mixed lifetimes

Cons:

- Garbage might survive a collection
- Small heap size for each generation
- More frequent collection